

Predict the price of the Uber ride from a given pickup point to the agreed drop-off location.

Perform following tasks:

1. Pre-process the dataset.
2. Identify outliers.
3. Check the correlation.
4. Implement linear regression and random forest regression models.
5. Evaluate the models and compare their respective scores like R2, RMSE, etc.

Dataset link:

<https://www.kaggle.com/datasets/yasserh/uber-fares-dataset>

```
In [1]: #Importing the required libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: #importing the dataset
df = pd.read_csv("C:\\Users\\Owner\\Desktop\\Machine Learning BE\\Practical\\Practi
df.head()
```

Out[2]:

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085



1) Pre-process the dataset

In [3]: `df.info() #To get the required information of the dataset`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   Unnamed: 0        200000 non-null   int64  
 1   key              200000 non-null   object  
 2   fare_amount       200000 non-null   float64 
 3   pickup_datetime   200000 non-null   object  
 4   pickup_longitude  200000 non-null   float64 
 5   pickup_latitude   200000 non-null   float64 
 6   dropoff_longitude 199999 non-null   float64 
 7   dropoff_latitude  199999 non-null   float64 
 8   passenger_count   200000 non-null   int64  
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
```

In [4]: `df.columns #TO get number of columns in the dataset`

Out[4]:

```
Index(['Unnamed: 0', 'key', 'fare_amount', 'pickup_datetime',
       'pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
       'dropoff_latitude', 'passenger_count'],
      dtype='object')
```

In [5]: `df = df.drop(['Unnamed: 0', 'key'], axis=1) #To drop unnamed column as it isn't re`

In [6]: `df.head()`

Out[6]:

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	-73.999512	40.7
1	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	-73.994710	40.7
2	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770	-73.962565	40.7
3	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	-73.965316	40.8
4	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085	-73.973082	40.7



In [7]: `df.shape #To get the total (Rows,Columns)`

Out[7]: `(200000, 7)`

In [8]: `df.dtypes #To get the type of each column`

Out[8]:

fare_amount	float64
pickup_datetime	object
pickup_longitude	float64
pickup_latitude	float64
dropoff_longitude	float64
dropoff_latitude	float64
passenger_count	int64
dtype:	object

In [9]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 7 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   fare_amount      200000 non-null  float64
 1   pickup_datetime  200000 non-null  object  
 2   pickup_longitude 200000 non-null  float64
 3   pickup_latitude   200000 non-null  float64
 4   dropoff_longitude 199999 non-null  float64
 5   dropoff_latitude   199999 non-null  float64
 6   passenger_count   200000 non-null  int64  
dtypes: float64(5), int64(1), object(1)
memory usage: 10.7+ MB
```

In [10]: `df.describe() #To get statistics of each columns`

Out[10]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
count	200000.000000	200000.000000	200000.000000	199999.000000	199999.000000	20
mean	11.359955	-72.527638	39.935885	-72.525292	39.923890	
std	9.901776	11.437787	7.720539	13.117408	6.794829	
min	-52.000000	-1340.648410	-74.015515	-3356.666300	-881.985513	
25%	6.000000	-73.992065	40.734796	-73.991407	40.733823	
50%	8.500000	-73.981823	40.752592	-73.980093	40.753042	
75%	12.500000	-73.967154	40.767158	-73.963658	40.768001	
max	499.000000	57.418457	1644.421482	1153.572603	872.697628	

In [11]: `df.isnull()`

Out[11]:

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False
...
199995	False	False	False	False	False	False	False
199996	False	False	False	False	False	False	False
199997	False	False	False	False	False	False	False
199998	False	False	False	False	False	False	False
199999	False	False	False	False	False	False	False

200000 rows × 7 columns

In [12]: `df.isnull().sum()`

Out[12]:

fare_amount	0
pickup_datetime	0
pickup_longitude	0
pickup_latitude	0
dropoff_longitude	1
dropoff_latitude	1
passenger_count	0
dtype:	int64

In [13]: `df.isnull().sum().sum()`

Out[13]: 2

2) Filling Missing values

```
In [14]: df.isnull().sum()
df['dropoff_latitude'].fillna(value=df['dropoff_latitude'].mean(),inplace = True)
df['dropoff_longitude'].fillna(value=df['dropoff_longitude'].median(),inplace = True)
df.isnull().sum()

Out[14]: fare_amount      0
pickup_datetime       0
pickup_longitude      0
pickup_latitude        0
dropoff_longitude      0
dropoff_latitude        0
passenger_count        0
dtype: int64
```

3) Column pickup_datetime is in wrong format (Object). Convert it to DateTime Format

```
In [15]: df.dtypes

Out[15]: fare_amount      float64
pickup_datetime       object
pickup_longitude      float64
pickup_latitude        float64
dropoff_longitude      float64
dropoff_latitude        float64
passenger_count        int64
dtype: object

In [16]: df.pickup_datetime = pd.to_datetime(df.pickup_datetime, errors='coerce',utc=True)

#df.pickup_datetime = pd.to_datetime(df.pickup_datetime)

In [17]: df.dtypes

Out[17]: fare_amount          float64
pickup_datetime    datetime64[ns, UTC]
pickup_longitude      float64
pickup_latitude        float64
dropoff_longitude      float64
dropoff_latitude        float64
passenger_count        int64
dtype: object
```

4) To segregate each time of date and time

```
In [18]: df.head()
```

Out[18]:

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	7.5	2015-05-07 19:52:06+00:00	-73.999817	40.738354	-73.999512	40.7
1	7.7	2009-07-17 20:04:56+00:00	-73.994355	40.728225	-73.994710	40.7
2	12.9	2009-08-24 21:45:00+00:00	-74.005043	40.740770	-73.962565	40.7
3	5.3	2009-06-26 08:22:21+00:00	-73.976124	40.790844	-73.965316	40.8
4	16.0	2014-08-28 17:47:00+00:00	-73.925023	40.744085	-73.973082	40.7

In [19]:

```
df = df.assign(hour = df.pickup_datetime.dt.hour, day = df.pickup_datetime.dt.day, month = df.pickup_datetime.dt.month)
df.head()
```

Out[19]:

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	7.5	2015-05-07 19:52:06+00:00	-73.999817	40.738354	-73.999512	40.7
1	7.7	2009-07-17 20:04:56+00:00	-73.994355	40.728225	-73.994710	40.7
2	12.9	2009-08-24 21:45:00+00:00	-74.005043	40.740770	-73.962565	40.7
3	5.3	2009-06-26 08:22:21+00:00	-73.976124	40.790844	-73.965316	40.8
4	16.0	2014-08-28 17:47:00+00:00	-73.925023	40.744085	-73.973082	40.7

In [20]:

```
# drop the column 'pickup_datetime' using drop()
# 'axis = 1' drops the specified column
df = df.drop(['pickup_datetime'], axis=1)
```

In [21]:

```
df.head()
```

Out[21]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	7.5	-73.999817	40.738354	-73.999512	40.723217	1
1	7.7	-73.994355	40.728225	-73.994710	40.750325	1
2	12.9	-74.005043	40.740770	-73.962565	40.772647	1
3	5.3	-73.976124	40.790844	-73.965316	40.803349	1
4	16.0	-73.925023	40.744085	-73.973082	40.761247	1

In [22]:

```
df.describe()
```

Out[22]:

	<code>fare_amount</code>	<code>pickup_longitude</code>	<code>pickup_latitude</code>	<code>dropoff_longitude</code>	<code>dropoff_latitude</code>	<code>passenger_count</code>
<code>count</code>	200000.000000	200000.000000	200000.000000	200000.000000	200000.000000	200000.000000
<code>mean</code>	11.359955	-72.527638	39.935885	-72.525299	39.923890	
<code>std</code>	9.901776	11.437787	7.720539	13.117375	6.794812	
<code>min</code>	-52.000000	-1340.648410	-74.015515	-3356.666300	-881.985513	
<code>25%</code>	6.000000	-73.992065	40.734796	-73.991407	40.733823	
<code>50%</code>	8.500000	-73.981823	40.752592	-73.980093	40.753042	
<code>75%</code>	12.500000	-73.967154	40.767158	-73.963659	40.768001	
<code>max</code>	499.000000	57.418457	1644.421482	1153.572603	872.697628	



```
In [23]: number_of_columns = len(df.columns)

In [24]: number_of_columns

Out[24]: 11

In [25]: #function to calculate the travel distance from the longitudes and latitudes
from math import *

def distance_formula(longitude1, latitude1, longitude2, latitude2):
    travel_dist = []

    for pos in range(len(longitude1)):
        lon1, lan1, lon2, lan2 = map(radians, [longitude1[pos], latitude1[pos], longitude2[pos], latitude2[pos]])
        dist_lon = lon2 - lon1
        dist_lan = lan2 - lan1

        a = sin(dist_lan/2)**2 + cos(lan1) * cos(lan2) * sin(dist_lon/2)**2

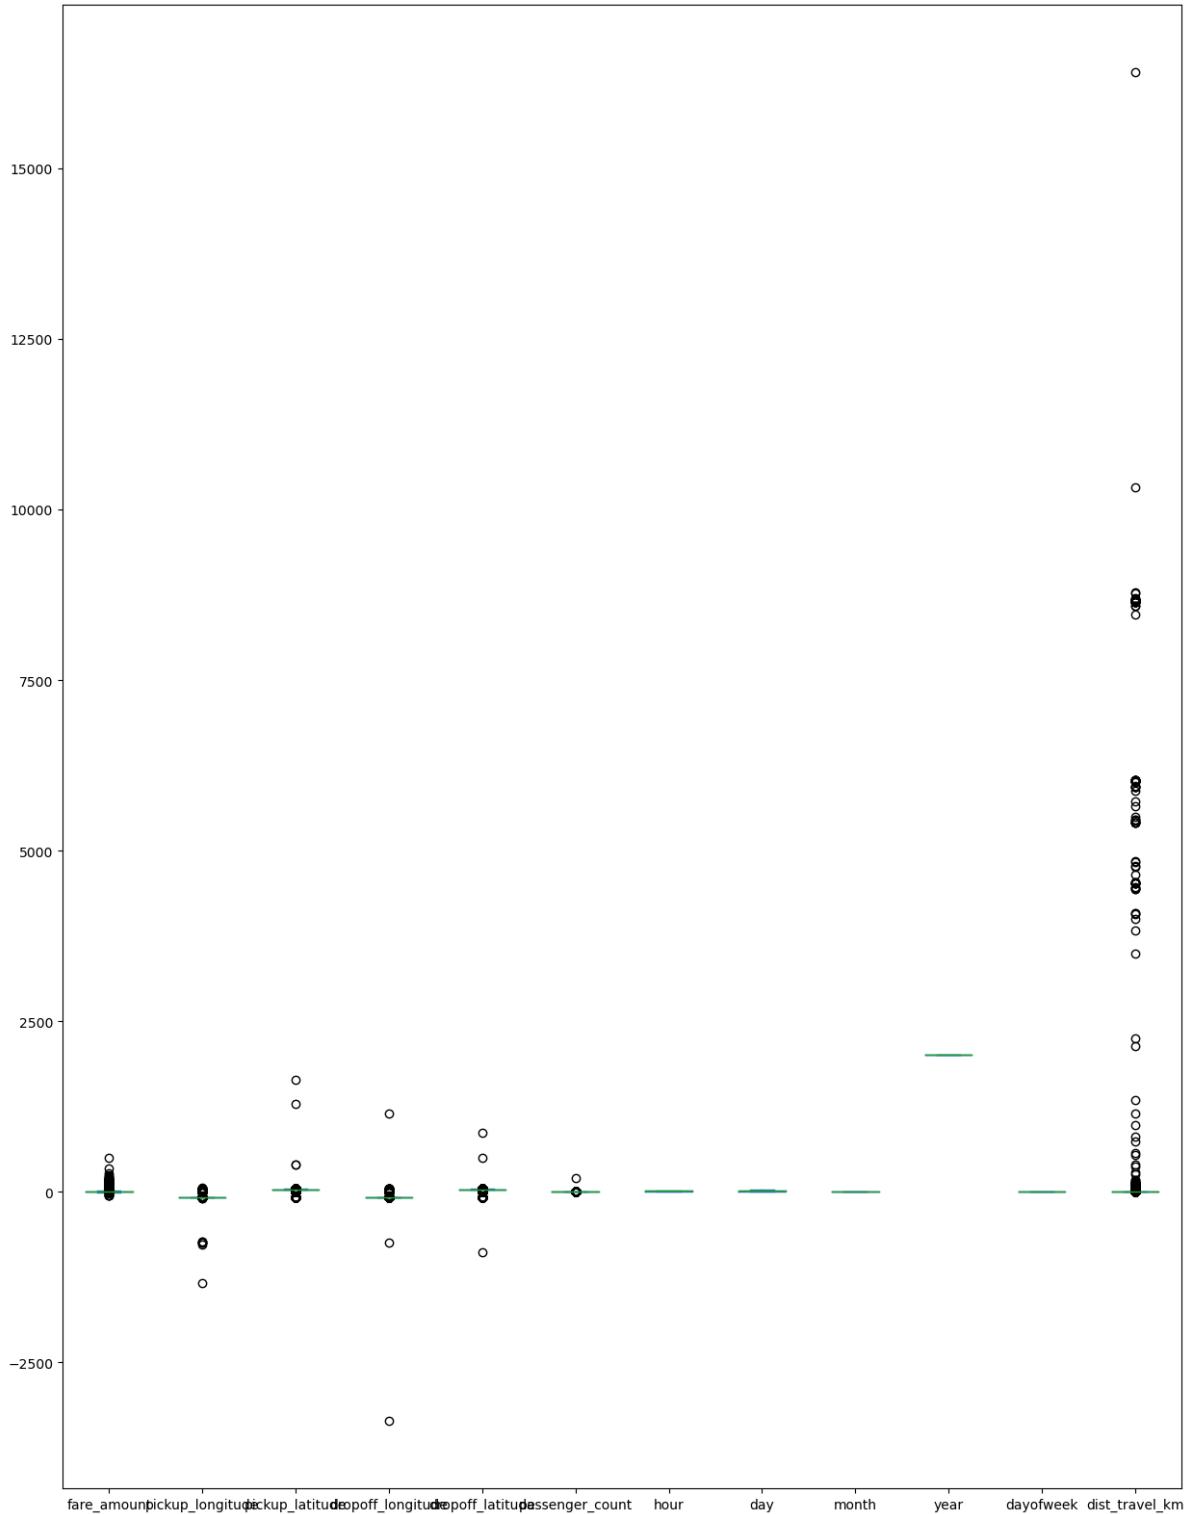
        #radius of earth = 6371
        c = 2 * asin(sqrt(a)) * 6371
        travel_dist.append(c)

    return travel_dist

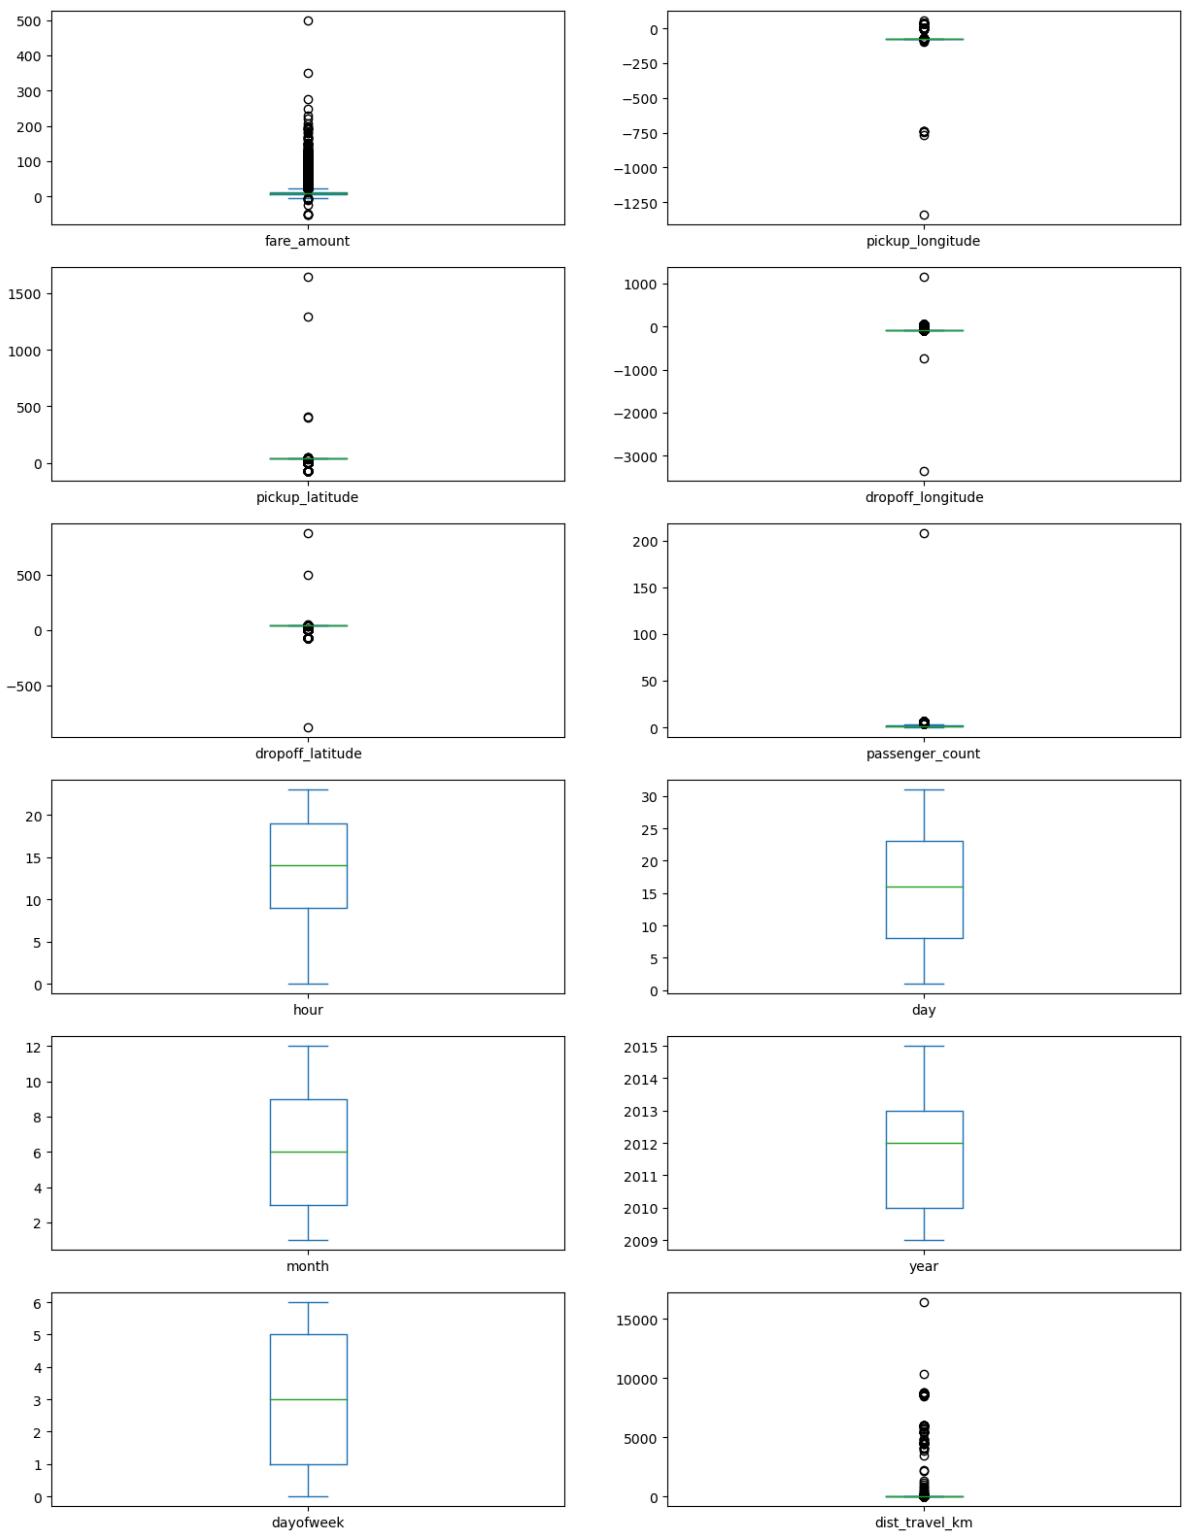
In [26]: df['dist_travel_km'] = distance_formula(df.pickup_longitude.to_numpy(), df.pickup_latitude.to_numpy())
```

5) Checking outliers and filling them

```
In [27]: df.plot(kind = "box", subplots = False, layout = (6,2), figsize=(15,20)) #Boxplot to check for outliers
plt.show()
```



```
In [28]: df.plot(kind = "box", subplots = True, layout = (6,2), figsize=(15,20)) #Boxplot to check the distribution of the data
plt.show()
```



In []:

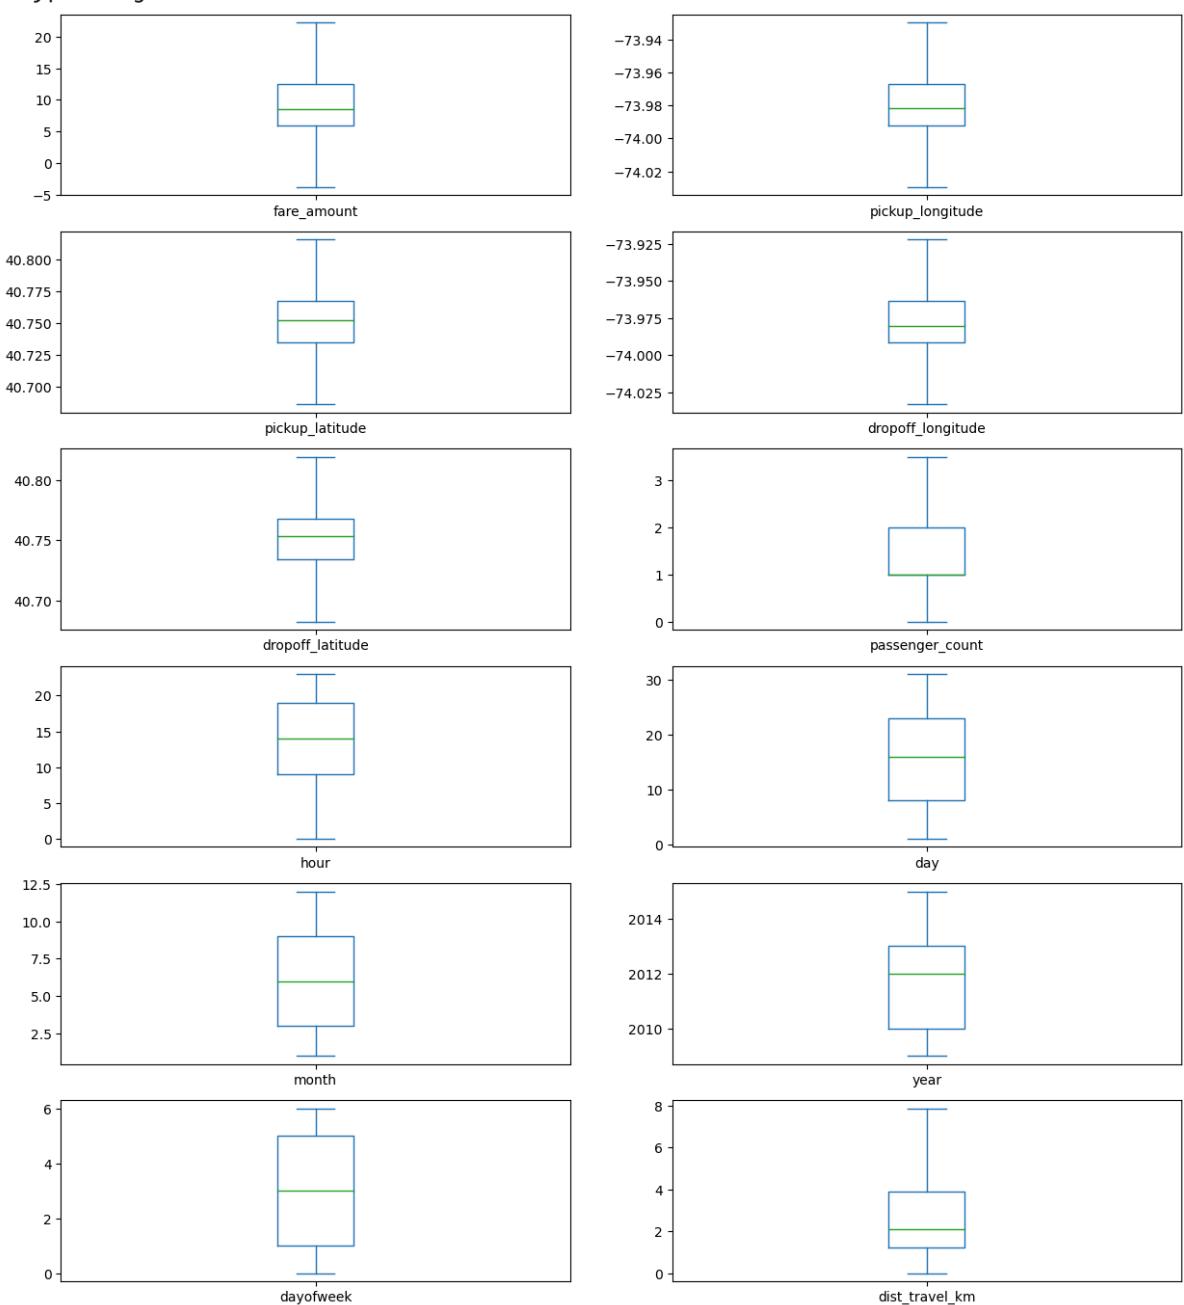
```

In [29]: def remove_outlier(df1 , col):
    Q1 = df1[col].quantile(0.25)
    Q3 = df1[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_whisker = Q1-1.5*IQR
    upper_whisker = Q3+1.5*IQR
    df1[col] = np.clip(df1[col] , lower_whisker , upper_whisker)
    return df1
def treat_outliers_all(df1 , col_list):
    for c in col_list:
        df1 = remove_outlier(df1 , c)
    return df1

```

```
df = treat_outliers_all(df, df.columns)
df.plot(kind = "box", subplots = True, layout = (7,2), figsize=(15,20)) #BoxPlot shows
```

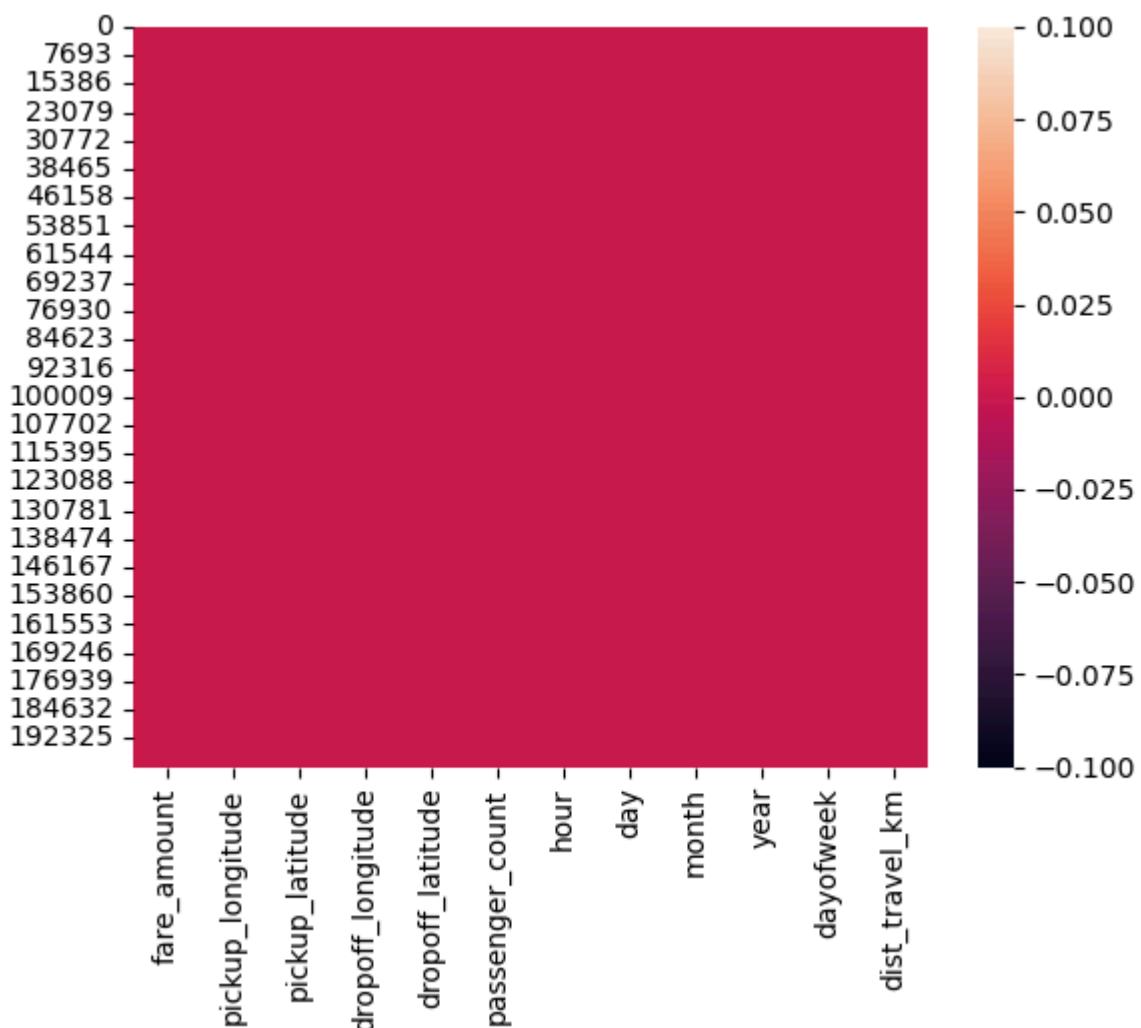
```
Out[29]: fare_amount      Axes(0.125,0.786098;0.352273x0.0939024)
pickup_longitude    Axes(0.547727,0.786098;0.352273x0.0939024)
pickup_latitude       Axes(0.125,0.673415;0.352273x0.0939024)
dropoff_longitude    Axes(0.547727,0.673415;0.352273x0.0939024)
dropoff_latitude      Axes(0.125,0.560732;0.352273x0.0939024)
passenger_count       Axes(0.547727,0.560732;0.352273x0.0939024)
hour                  Axes(0.125,0.448049;0.352273x0.0939024)
day                   Axes(0.547727,0.448049;0.352273x0.0939024)
month                 Axes(0.125,0.335366;0.352273x0.0939024)
year                  Axes(0.547727,0.335366;0.352273x0.0939024)
dayofweek              Axes(0.125,0.222683;0.352273x0.0939024)
dist_travel_km         Axes(0.547727,0.222683;0.352273x0.0939024)
dtype: object
```



```
In [30]: #Finding incorrect latitude (Less than or greater than 90) and longitude (greater than 180)
incorrect_coordinates = df.loc[(df.pickup_latitude > 90) | (df.pickup_latitude < -90) |
                               (df.dropoff_latitude > 90) | (df.dropoff_latitude < -90) |
                               (df.pickup_longitude > 180) | (df.pickup_longitude < -180) |
                               (df.dropoff_longitude > 90) | (df.dropoff_longitude < -90) ]
df.drop(incorrect_coordinates, inplace = True, errors = 'ignore')
df.head()
```

```
df.isnull().sum()  
sns.heatmap(df.isnull()) #Free for null values
```

Out[30]: <Axes: >



In [31]: incorrect_coordinates

Out[31]: fare_amount pickup_longitude pickup_latitude dropoff_longitude dropoff_latitude passenger_c

In [32]: #Uber doesn't travel over 130 kms so minimize the distance
df = df.loc[(df.dist_travel_km >= 1) | (df.dist_travel_km <= 130)]
print("Remaining observations in the dataset:", df.shape)

Remaining observations in the dataset: (200000, 12)

7)Check the correlation.

```
#Function to find the correlation  
corr = df.corr()  
corr
```

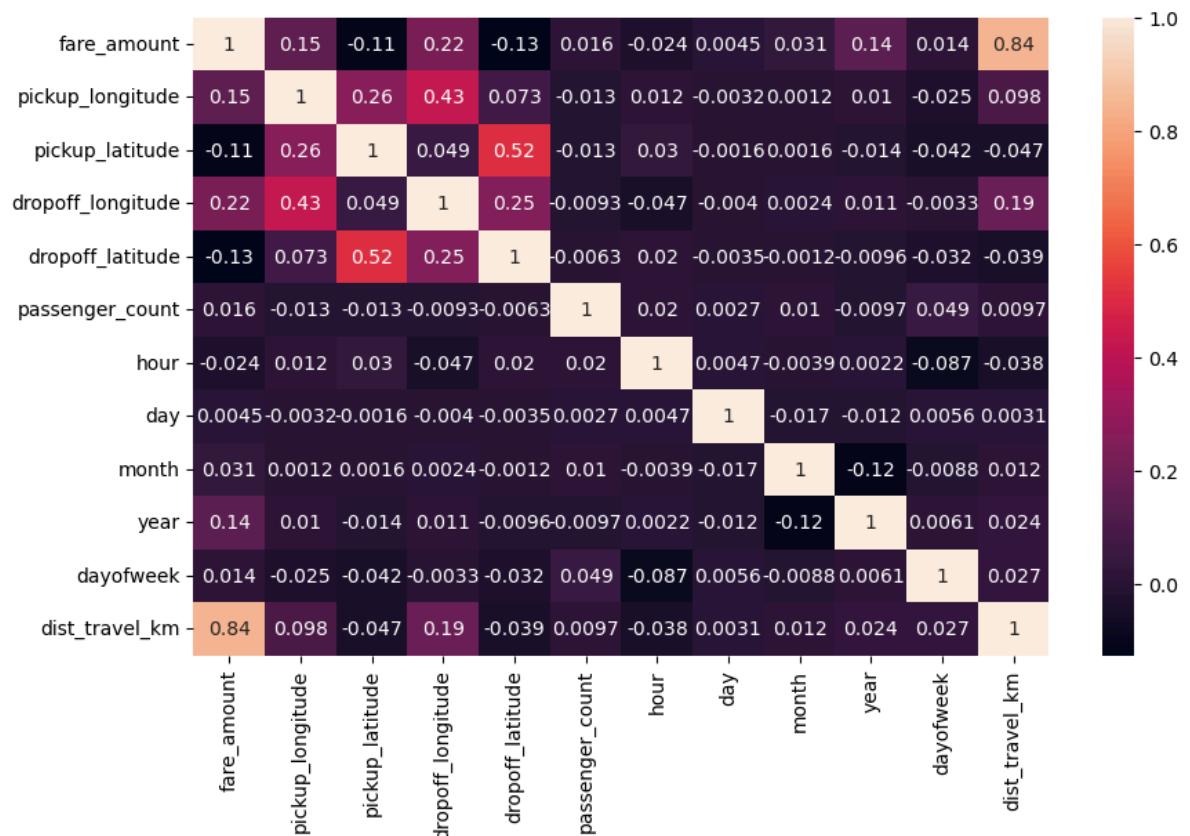
Out[33]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
fare_amount	1.000000	0.154069	-0.110842	0.218675	-0.125
pickup_longitude	0.154069	1.000000	0.259497	0.425619	0.073
pickup_latitude	-0.110842	0.259497	1.000000	0.048889	0.515
dropoff_longitude	0.218675	0.425619	0.048889	1.000000	0.245
dropoff_latitude	-0.125898	0.073290	0.515714	0.245667	1.000
passenger_count	0.015778	-0.013213	-0.012889	-0.009303	-0.006
hour	-0.023623	0.011579	0.029681	-0.046558	0.019
day	0.004534	-0.003204	-0.001553	-0.004007	-0.003
month	0.030817	0.001169	0.001562	0.002391	-0.001
year	0.141277	0.010198	-0.014243	0.011346	-0.009
dayofweek	0.013652	-0.024652	-0.042310	-0.003336	-0.031
dist_travel_km	0.844374	0.098094	-0.046812	0.186531	-0.038

In [34]:

```
fig, axis = plt.subplots(figsize = (10,6))
sns.heatmap(df.corr(), annot = True) #Correlation Heatmap (Light values means highly
```

Out[34]:



6)Dividing the dataset into feature and target value

```
In [35]: df.columns
```

```
Out[35]: Index(['fare_amount', 'pickup_longitude', 'pickup_latitude',
   'dropoff_longitude', 'dropoff_latitude', 'passenger_count', 'hour',
   'day', 'month', 'year', 'dayofweek', 'dist_travel_km'],
  dtype='object')
```

```
In [36]: df.dtypes
```

```
Out[36]: fare_amount      float64
pickup_longitude    float64
pickup_latitude      float64
dropoff_longitude    float64
dropoff_latitude      float64
passenger_count      float64
hour                  int64
day                   int64
month                 int64
year                  int64
dayofweek              int64
dist_travel_km      float64
dtype: object
```

```
In [37]: x = df[['hour', 'day', 'month', 'year', 'dayofweek', 'dist_travel_km']]
y = df['fare_amount']
```

7) Scaling and Dividing the dataset into training and testing dataset

```
In [38]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(x)
```

```
In [39]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X_scaled,y,test_size = 0.25)
```

```
In [40]: X_train
```

```
Out[40]: array([[ 0.69198905, -1.69264582,  1.66279263,  0.13874222,  1.00238004,
   0.81391428],
   [ 1.45938811, -0.88687998, -0.08194296,  0.13874222,  0.48875385,
   1.22698408],
   [-2.07064754,  0.26421407,  0.2088463 ,  0.13874222, -0.53849854,
  -1.03599597],
   ...,
   [-0.68932925,  0.9548705 , -1.24510003, -1.4772954 , -1.05212474,
  2.25341213],
   [ 0.38502943,  0.26421407, -1.24510003, -0.9386162 , -0.02487235,
  -0.86046649],
   [ 1.30590829, -1.11709879,  1.37200337, -0.9386162 ,  1.00238004,
  -0.9189509 ]])
```

```
In [41]: y_train
```

```
Out[41]: 32942    20.50
          175963   14.10
          15734    3.30
          88189    5.70
          151034   17.30
          ...
          77880    8.90
          120656   5.30
          105690   22.25
          56691    3.70
          21554    10.50
Name: fare_amount, Length: 150000, dtype: float64
```

```
In [42]: len(y)
```

```
Out[42]: 200000
```

```
In [43]: len(y_train)
```

```
Out[43]: 150000
```

```
In [44]: len(y_test)
```

```
Out[44]: 50000
```

8)Linear Regression

```
In [45]: from sklearn.linear_model import LinearRegression
regression = LinearRegression()
```

```
In [46]: regression.fit(X_train,y_train)
```

```
Out[46]: ▾ LinearRegression
LinearRegression()
```

```
In [47]: regression.coef_ #To find the linear coefficient
```

```
Out[47]: array([ 0.03649165,  0.02062     ,  0.19146925,  0.68237769, -0.04565736,
        4.57557551])
```

```
In [48]: regression.intercept_ #To find the linear intercept
```

```
Out[48]: 10.076692713222002
```

```
In [49]: for i in range(0,len(regression.coef_)):
print("Theta",i,"=",regression.coef_[i])
```

```
Theta 0 = 0.036491647419307924
Theta 1 = 0.02062000382452531
Theta 2 = 0.19146924507828567
Theta 3 = 0.6823776928080336
Theta 4 = -0.04565735884232406
Theta 5 = 4.575575509222712
```

```
In [50]: y_prediction = regression.predict(X_test) #To predict the target values
print(y_prediction)
```

```
[ 8.35461279  8.30441313  7.68471507 ...  7.56411998 20.54416486  
 5.65656427]
```

```
In [51]: y_test
```

```
Out[51]:
```

110440	7.00
170545	6.50
115711	10.50
158343	8.90
115299	9.00
	...
147278	13.50
89654	22.25
58881	6.90
183886	22.25
49177	5.30

```
Name: fare_amount, Length: 50000, dtype: float64
```

```
In [52]: comparison=pd.DataFrame({"Actual Value":y_test, "Predicted Value":y_prediction })
```

```
In [53]: comparison
```

```
Out[53]:
```

	Actual Value	Predicted Value
110440	7.00	8.354613
170545	6.50	8.304413
115711	10.50	7.684715
158343	8.90	9.254395
115299	9.00	6.520236
...
147278	13.50	14.636705
89654	22.25	20.855361
58881	6.90	7.564120
183886	22.25	20.544165
49177	5.30	5.656564

50000 rows × 2 columns

```
In [54]: comparison.reset_index()
```

Out[54]:

	index	Actual Value	Predicted Value
0	110440	7.00	8.354613
1	170545	6.50	8.304413
2	115711	10.50	7.684715
3	158343	8.90	9.254395
4	115299	9.00	6.520236
...
49995	147278	13.50	14.636705
49996	89654	22.25	20.855361
49997	58881	6.90	7.564120
49998	183886	22.25	20.544165
49999	49177	5.30	5.656564

50000 rows × 3 columns

In [55]:

```
comparison.reset_index().drop(["index"], axis=1)
```

Out[55]:

	Actual Value	Predicted Value
0	7.00	8.354613
1	6.50	8.304413
2	10.50	7.684715
3	8.90	9.254395
4	9.00	6.520236
...
49995	13.50	14.636705
49996	22.25	20.855361
49997	6.90	7.564120
49998	22.25	20.544165
49999	5.30	5.656564

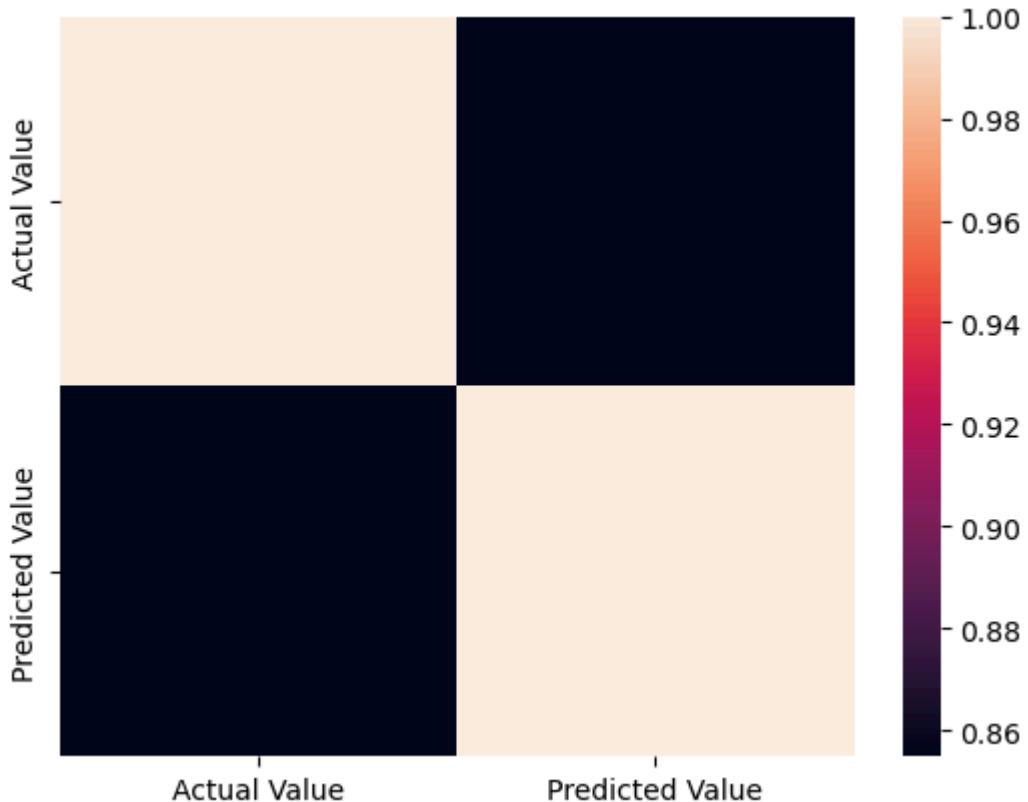
50000 rows × 2 columns

In [56]:

```
sns.heatmap(comparison.corr())
```

Out[56]:

```
<Axes: >
```



9) Metrics Evaluation using R2, Mean Squared Error, Root Mean Squared Error

```
In [57]: from sklearn.metrics import r2_score
r2_score(y_test,y_prediction)
```

```
Out[57]: 0.730770273098684
```

```
In [58]: from sklearn.metrics import mean_squared_error
MSE = mean_squared_error(y_test,y_prediction)
MSE
```

```
Out[58]: 7.976966086868074
```

```
In [59]: RMSE = np.sqrt(MSE)
RMSE
```

```
Out[59]: 2.8243523305119127
```

10) Random Forest Regression

```
In [65]: from sklearn.ensemble import RandomForestRegressor
rf = RandomForestRegressor(n_estimators=100) #Here n_estimators means number of trees
rf.fit(X_train,y_train)
y_pred = rf.predict(X_test)
y_pred
```

```
Out[65]: array([ 8.1 ,  8.064,  8.4 , ...,  7.196, 17.36 ,  5.559])
```

```
In [66]: from sklearn.metrics import r2_score  
r2_score(y_test,y_prediction)
```

```
Out[66]: 0.730770273098684
```

```
In [67]: from sklearn.metrics import mean_squared_error  
MSE = mean_squared_error(y_test,y_prediction)  
MSE
```

```
Out[67]: 7.976966086868074
```

```
In [68]: RMSE = np.sqrt(MSE)  
RMSE
```

```
Out[68]: 2.8243523305119127
```

```
In [ ]:
```

```
In [ ]:
```