

```
[2]: from sklearn.neural_network import MLPClassifier
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from imblearn.over_sampling import RandomOverSampler
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
```

```
[3]: df = pd.read_csv('Churn_Modelling.csv')
```

```
[4]: df
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	E
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88	
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57	
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	
...	
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.00	2	1	0	96270.64	
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	1	1	101699.77	
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	1	42085.58	

```
[5]: df.isnull().sum()
```

```
[5]: RowNumber      0
CustomerId      0
Surname        0
CreditScore     0
Geography       0
Gender          0
Age             0
Tenure          0
Balance         0
NumOfProducts   0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
dtype: int64
```

```
[16]: df['Exited'].value_counts()
```

```
[16]: Exited
0    7963
1    2037
Name: count, dtype: int64
```

```
[6]: X = df.drop(['RowNumber', 'CustomerId', 'Surname', 'Exited'], axis=1)
y = df['Exited']
```

```
[7]: X = pd.get_dummies(X, drop_first=True)
```

```
[8]: ros = RandomOverSampler(random_state=42)
X_resampled, y_resampled = ros.fit_resample(X, y)

print("Class distribution after balancing:")
print(y_resampled.value_counts())
```

```
Class distribution after balancing:
Exited
1    7963
0    7963
Name: count, dtype: int64
```

```
*[9]: X_train, X_test, y_train, y_test = train_test_split(X_resampled, y_resampled, test_size=0.2, random_state=42)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
[22]: X_train_df = pd.DataFrame(X_train)
X_test_df = pd.DataFrame(X_test)
```

```
[24]: X_train_df.head()
[24]:
   0      1      2      3      4      5      6      7      8      9      10
0  1.465814  0.079328  1.719402  0.511082  0.729063 -1.530958  1.082515  0.692593  1.509728 -0.552686 -1.03061
1  1.035041  0.079328 -1.364387  0.773670  0.729063 -1.530958  1.082515  0.755225 -0.662371 -0.552686  0.97030
2  0.522216 -0.109329  0.691472 -1.340431 -0.758963 -1.530958  1.082515  0.918049 -0.662371 -0.552686  0.97030
3  0.101700  1.305597  1.034116  0.603456 -0.758963 -1.530958  1.082515  0.335004  1.509728 -0.552686 -1.03061
4  0.081187  0.456641  0.006186  0.506923  0.729063  0.653186  1.082515  0.673453 -0.662371 -0.552686 -1.03061

[25]: X_test_df.head()
[25]:
   0      1      2      3      4      5      6      7      8      9      10
0 -1.159848 -1.146942  0.006186 -1.340431  0.729063 -1.530958  1.082515 -0.828594 -0.662371 -0.552686 -1.03061
1 -0.431637 -0.015001 -0.336457  0.820811 -0.758963  0.653186 -0.923775  0.980827  1.509728 -0.552686 -1.03061
2  0.327343 -0.958285  1.376759  0.611566  0.729063  0.653186  1.082515  0.015894 -0.662371  1.809346 -1.03061
3  0.450421 -0.486643 -0.336457 -1.340431 -0.758963  0.653186 -0.923775  1.488334 -0.662371 -0.552686 -1.03061
4  0.388882 -0.203658 -0.336457  1.229063  0.729063  0.653186 -0.923775  1.091454  1.509728 -0.552686 -1.03061

[18]: y_train.value_counts()
[18]:
Exited
1    6410
0    6330
Name: count, dtype: int64

[17]: y_test.value_counts()
[17]:
Exited
0    1633
1    1553
Name: count, dtype: int64

[10]: mlp_model = MLPClassifier(hidden_layer_sizes=(64, 32), activation='relu', solver='adam', max_iter=200, random_state=42)
mlp_model.fit(X_train, y_train)
y_pred = mlp_model.predict(X_test)

C:\Users\shreyash\AppData\Roaming\Python\Python311\site-packages\sklearn\neural_network\_multilayer_perceptron.py:691: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
  warnings.warn(
[11]: accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy * 100:.2f}%")
Accuracy: 84.59%


[12]: print("\nClassification Report:")
print(classification_report(y_test, y_pred))

Classification Report:
              precision    recall  f1-score   support

             0       0.90      0.79      0.84     1633
             1       0.80      0.91      0.85     1553

    accuracy                           0.85     3186
   macro avg       0.85      0.85      0.85     3186
weighted avg       0.85      0.85      0.85     3186

[13]: conf_matrix = confusion_matrix(y_test, y_pred)
print('Confusion Matrix:')
print(conf_matrix)
Confusion Matrix:
[[1289 344]
 [ 147 1406]]


[14]: plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix, annot=True, fmt='g', cmap='Blues', cbar=False)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()



```